

# Introduction to Machine Learning - Homework 1

Savinay Shukla

Q1 ) Below are the snapshots of the polyreg function observed with different values of  $d$  :




Figure 1.1 - polyreg plot with  $d = 1$ .




Figure 1.2 - polyreg plot with  $d = 2$ .




Figure 1.3 Polyreg graph with  $d = 3$




Figure 1.4 polyreg graph with  $d = 4$ .

We can observe that the polynomial function begins to closely approximate the data distribution as we increase the value of “ $d$ ”.

Below is the use of a cross-validation plot with various values of “ $d$ ” to find the best order of polynomial that would fit this data.




Figure 1.5 Cross-validation

As we can see from the plot above, the minimum error for the testing data is coming at the **degree = 6**.

For degrees  $> 6$ , we find that the test error starts to increase rapidly, and begins to overfit the model. So according to the data given, **we can say that the best order of the polynomial = 6, i.e.  $d = 6$** . We can confirm this by plotting the curve against the data distribution. The curve will fit almost perfectly with the data at  $d = 6$ .

Below is the graph plot of with  $d = 6$ .




Figure 1.6 The polyreg function plot with optimal value of  $d = 6$ .

### Source Code (MATLAB) :

```
clear all;
load("problem1.mat");
indx = crossvalind("Kfold",length(x), 2);
xtrain = x(indx == 1);
xtest = x(indx == 2);
ytrain = y(indx == 1);
ytest = y(indx == 2);
trainError = [];
testError = [];
iterations = [];
for i=1:40
    [err,model,errT] = polyreg(xtrain,ytrain,i,xtest,ytest);
    trainError = [trainError, err];
    testError = [testError, errT];
    iterations = [iterations,i];
end
figure
plot(iterations,trainError,iterations,testError, 'LineWidth',2.0)
xlabel("Polynomial Degree");
ylabel("Error");
legend("Train Error", "Test Error");
title("Cross-Validation Performance");
```

---

### Polyreg.m

```
function [err,model,errT] = polyreg(x,y,D,xT,yT)
%
% Finds a D-1 order polynomial fit to the data
%
%     function [err,model,errT] = polyreg(x,y,D,xT,yT)
%
% x = vector of input scalars for training
% y = vector of output scalars for training
% D = the order plus one of the polynomial being fit
% xT = vector of input scalars for testing
% yT = vector of output scalars for testing
% err = average squared loss on training
% model = vector of polynomial parameter coefficients
% errT = average squared loss on testing
%
% Example Usage:
%
% x = 3*(rand(50,1)-0.5);
% y = x.*x.*x-x+rand(size(x));
% [err,model] = polyreg(x,y,4);
%
xx = zeros(length(x),D);
for i=1:D
    xx(:,i) = x.^(D-i);
end
model = pinv(xx)*y;
err = (1/(2*length(x)))*sum((y-xx*model).^2);
if (nargin==5)
    xxT = zeros(length(xT),D);
    for i=1:D
        xxT(:,i) = xT.^(D-i);
    end
    errT = (1/(2*length(xT)))*sum((yT-xxT*model).^2);
end
q = (min(x):(max(x)/300):max(x))';
qq = zeros(length(q),D);
for i=1:D
    qq(:,i) = q.^(D-i);
end
clf
    plot(x,y,'X');
hold on
if (nargin==5)
    plot(xT,yT,'cO');
end
    plot(q,qq*model,'r')
```

Q2)




Figure 2.1 Error vs Penalty Parameter.

We can see that as we increase the penalty parameter – *lambda*, we observe a decrease in testing error. We also observe an approximate converging trend of error curves ( both test error curve and train error curve ) with the above graph plot.

Additionally, we see that after an approximate value of  $\lambda = 750$ , there is not much improvement in test error. However, this behavior is believed to be dependent on how the test-train data was initially split.

Below is an example of the test-train error curve observed after running the source code multiple times.




Figure 2.2 Converging Error vs Penalty Parameter

In the above example, we observe a convergence of test errors and train errors **at  $\lambda = 1000$** .

## Source Code(MATLAB)

### Driver code:

```
clear all;
load("problem2.mat");
clc;
indx = crossvalind("Kfold",length(x), 2);
x_train = x(indx == 1,:);
x_test = x(indx == 2,:);
y_train = y(indx == 1);
y_test = y(indx == 2);
trainErrors = [];
testErrors = [];
iterations = [];
%Creating a lambda Set from 0 - 1000 lambda values.
lambdaSet = 0:0.2:1000;
for i=1:length(lambdaSet)
    [err,theta,errT] = regularizedReg(x_train,y_train,i,x_test,y_test);
    trainErrors = [trainErrors, err];
    testErrors = [testErrors, errT];
    iterations = [iterations,i];
end
figure
hold on;
plot(iterations,trainErrors,iterations,testErrors,'LineWidth',2.0);
xlabel('Lambda');
ylabel('Error');
legend('Train Error', 'Test Error');
hold off;
```

---

### regularizedReg.m

```
function [trainErr,theta,testErr] = regularizedReg(x,y,lambda,xT,yT)
% This method returns "trainErr" - regularized error during training.
%
% "theta" - vector of multivariate coefficients.
%
% "testErr" - regularized error during testing.
% This method takes the following parameters:
%
% "x" - training feature vector;
%
% "y" - training target vector;
%
% "xT" - testing feature vector;
%
% "yT" - testing target vector;
%
% lambda - penalty parameter;
%
%
theta = (x' * x + lambda * eye(size(x'* x))) \ x' * y;
trainErr = (1 / (2 * length(x))) * sum((y - x*theta).^2);
trainErr = trainErr + (lambda/(2*(length(x)))) * (theta' * theta);
```



```

if(nargin == 5)
    testErr = (1 / (2 * length(xT))) * sum((yT - xT*theta).^2);
    testErr = testErr + (lambda/(2*(length(xT)))) * (theta' * theta);
end
end

```

Q3)

PROOF FOR :  $g(-z) = 1 - g(z)$

$$\begin{aligned}
 g(z) &= \frac{1}{1+e^{-z}} \\
 &= \frac{1}{1+\frac{1}{e^z}} = \frac{e^z}{1+e^z}
 \end{aligned}$$

$$\begin{aligned}
 1 - g(z) &= 1 - \frac{e^z}{1+e^z} \\
 &= \frac{1+e^z - e^z}{1+e^z} = \frac{1}{1+e^z} \quad \hookrightarrow \textcircled{1}
 \end{aligned}$$

---


$$\begin{aligned}
 g(-z) &= \frac{1}{1+e^{-(-z)}} \\
 &= \frac{1}{1+e^z} \quad \text{--- } \textcircled{2}
 \end{aligned}$$

Since  $\textcircled{1} = \textcircled{2}$

Hence:  $g(-z) = 1 - g(z)$

PROOF FOR:  $g^{-1}(y) = \ln\left(\frac{y}{1-y}\right)$

Given that  $g(z) = y$ , we need to prove that  $g^{-1}(y) = z$   
 $= \ln\left(\frac{y}{1-y}\right) \quad \text{--- (1)}$

Solving for  $\ln\left(\frac{y}{1-y}\right)$  :-

$$\begin{aligned}\ln\left(\frac{y}{1-y}\right) &= \ln(y) - \ln(1-y) \\ &= \ln\left(\frac{1}{1+e^{-z}}\right) - \ln\left(1 - \frac{1}{1+e^{-z}}\right) \\ &\quad \left\{ \text{Because } y = \frac{1}{1+e^{-z}} \right. \\ &= \ln\left(\frac{1}{1+e^{-z}}\right) - \ln\left(\frac{e^{-z}}{1+e^{-z}}\right) \\ &= \ln(1) - \ln(1+e^{-z}) - \ln(e^{-z}) \\ &\quad \quad \quad + \ln(1+e^{-z}) \\ &= \ln(1) - \ln(e^{-z}) \\ &= -(-z) \\ &= z\end{aligned}$$

$$\ln\left(\frac{y}{1-y}\right) = z = g^{-1}(y)$$

Hence Proved!

Q4)

The gradient descent function is as follows :

$$\theta^1 = \theta^0 - \eta \nabla_{\theta} R_{emp} \Big|_{\theta^0}, \quad t = 1$$

So before we optimize our theta vector with the gradient descent algorithm, we need to find the derivative of our empirical risk function.

Given risk function is :

$$R_{emp}(\boldsymbol{\theta}) = \frac{1}{N} \sum_{i=1}^N (y_i - 1) \log(1 - f(\mathbf{x}_i; \boldsymbol{\theta})) - y_i \log(f(\mathbf{x}_i; \boldsymbol{\theta})).$$

Additionally the given classification function is:

$$f(\mathbf{x}; \boldsymbol{\theta}) = \left(1 + \exp(-\boldsymbol{\theta}^{\top} \mathbf{x})\right)^{-1}$$

Below is described the way to obtain the derivative (gradients) for the given risk function:

$$\text{Let } \theta^T u = t$$

So after putting the value of  $\theta^T u$  in  $f(x; \theta)$  we get:-

$$f(t) = \frac{1}{1 + e^{-t}}$$

$$\text{Let } f(t) = s$$

We can derive  $\text{Remp}(\theta)$  as follows.

$$\frac{d \text{Remp}(\theta)}{d\theta} = \frac{d \text{Remp}(\theta)}{ds} \times \frac{ds}{dt} \times \frac{dt}{d\theta} \quad (\text{Product of partial derivatives})$$

Solving each derivative separately:-

$$\begin{aligned} \frac{d \text{Remp}(\theta)}{ds} &= \frac{d}{ds} \left( (y-1) \log(1-s) - y \log(s) \right) \\ &= \frac{1-y}{1-s} - \frac{y}{s} \\ &= \frac{s - ys - y + ys}{s(1-s)} \\ &= \frac{s-y}{s(1-s)} \quad \text{--- (1)} \end{aligned}$$

$$\begin{aligned} \frac{ds}{dt} &= \frac{d}{dt} \left( \frac{1}{1+e^{-t}} \right) \\ &= \frac{d}{dt} \left( \frac{e^t}{1+e^t} \right) \\ &= \frac{e^t(1+e^t) - e^{2t}}{(1+e^t)^2} \end{aligned}$$

$$= \frac{e^t + e^{2t} - e^{2t}}{(1+e^t)^2}$$

$$= \frac{e^t}{(1+e^t)^2} = \frac{e^t}{(1+e^t)} \times \frac{1}{(1+e^t)}$$

Since  $\frac{e^t}{1+e^t} = s$   
 $\hookrightarrow s = s \times (1-s)$

$$\therefore \frac{ds}{dt} = s(1-s) \quad \text{--- (2)}$$

$$\frac{dt}{d\theta} = \frac{d}{d\theta} (\theta^T x)$$

$$= x \quad \text{--- (3)}$$

$$\frac{d \text{Kemp}(\theta)}{d\theta} = \textcircled{1} \times \textcircled{2} \times \textcircled{3}$$

$$= \left( \frac{s-y}{s(1-s)} \right) \times (s(1-s)) \times (x)$$

$$= (s-y)x$$

Substituting back the values of  $s$  and  $t$ ,

$$\frac{d \text{Kemp}(\theta)}{d\theta} = (f(x; \theta) - y_i) x_i$$

$$\therefore \nabla \text{Kemp}(\theta) = \frac{1}{N} \sum (f(x; \theta) - y_i) x_i$$

$\hookrightarrow$  To be used to find gradients.

Here are the graph plots for different values of **n** and **e**:




Figure 4.1 Decision boundary with  $n = 0.1$  and  $\text{epsilon} = 0.001$

The above graph is the decision boundary obtained when the gradient descent algorithm is applied with  **$n = 0.1$**  and  **$\text{epsilon} = 0.001$** . The decision boundary is not accurately classifying the data as per these values of **n** and **e**.




Figure 4.2 Risk and Classification Error versus Iterations.

The above is a risk plot and binary classification error plot for the same value of  $n=0.01$  and  $e=0.001$ . Number of iterations recorded = approx 4000.




Figure 4.3 Decision Boundary Plot for  $n = 0.5$  and  $\epsilon = 0.001$

The above is the decision boundary plot for  $n = 0.5$  and  $\epsilon = 0.001$ . We observe better accuracy with these values of  $n$  and  $\epsilon$  than before. Recorded iterations were around 4000.






Figure 4.4 Risk and Classification Error versus Iterations for  $n = 0.5$  and  $\epsilon = 0.001$

Below is the decision boundary plot for  $n=1$  and  $\epsilon = 0.001$ :




Figure 4.5 Decision Boundary Plot for  $n = 1$  and  $\epsilon = 0.001$




Figure 4.6 Risk and Classification Error vs Iterations for  $n = 1$  and  $\epsilon = 0.001$

We have observed that with  $n = 1$  and  $\epsilon = 0.001$ , we get the **minimum classification error of 1%** and lowest risk at almost 3000 iterations. In addition, the decision boundary seems to classify the data points with most accuracy using these values of  $n$  and  $e$ .

Source :

**Driver code:**

```
%% Load DataSet
clear all;
load("dataset4.mat");
%% Parameters
stepSize =1;
epsilon = 0.001;
upperBoundLimit = 8000; %Set the maximum number of iterations
%% Training
[row, col] = size(X);
initialTheta = rand(size(X, 2),1); %random initialization of theta;
currTheta = initialTheta + epsilon; %creating a currentTheta vector which has greater
values than initialTheta;
iterations = [];
costPerIteration = [];
errors=[];
i = 1;
while norm(currTheta - initialTheta) >= epsilon

    %an upperBound limit on number of iterations
    if i > upperBoundLimit
        break;
    end
    [risk, gradients] = computeCost(X, Y, initialTheta);
    prediction = 1 ./ ( 1 + exp (-X * initialTheta));
    prediction( prediction >= 0.5) = 1;
    prediction( prediction < 0.5) = 0;
    %calculate the error = (sum of all predictions != Y)/length(Y)
    err = sum(prediction~=Y)/length(Y);
    errors = [errors, err];
    costPerIteration = [costPerIteration, risk];
    %Append the iteration count for plotting.
    iterations = [iterations, i];
    i = i + 1;
    %Update the theta vector;
    currTheta = initialTheta;
    initialTheta = initialTheta - stepSize * gradients;
end
%For Plotting Risk and Classification Error graphs.
figure
plot(iterations, costPerIteration,iterations, errors,'LineWidth',2.0);
legend('Risk','Classification Error');
ylabel('Error');
xlabel('Iterations');
figure
%For Plotting Decision Boundary
```

```

x = 0:0.001:1;
y = (-initialTheta(3) - initialTheta(1).*x)/initialTheta(2);
plot(x, y, "b"); hold on;
plot(X(:, 1), X(:, 2), "r");
title("Decision Boundary")

```

### computeCost.m

```

%% Function for computing risk and gradients.
function [cost, gradients] = computeCost(x, y, theta)
h = sigmoid( x * theta);
cost = - (1 / length(x)) * sum ( y .* log(h) + (1-y).*log(1-h));
gradients = zeros(length(theta), 1);
for i = 1:size(gradients)
    gradients(i) = (1 / length(x)) * sum((h - y)' * x(:,i));
end
end

```

### sigmoid.m

```

%% Function for Logistic Squashing functions.
function [sig] = sigmoid(x)
sig = 1 ./ (1 + exp(-x));
end

```